

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

## Load California Housing Dataset

```
In [2]: data = fetch_california_housing()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['Target'] = data.target
```

## Display basic info

In [3]:

```
print(df.info())
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	MedInc	20640 non-null	float64
1	HouseAge	20640 non-null	float64
2	AveRooms	20640 non-null	float64
3	AveBedrms	20640 non-null	float64
4	Population	20640 non-null	float64
5	AveOccup	20640 non-null	float64
6	Latitude	20640 non-null	float64
7	Longitude	20640 non-null	float64
8	Target	20640 non-null	float64

```
dtypes: float64(9)
```

```
memory usage: 1.4 MB
```

```
None
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744
std	1.899822	12.585558	2.474173	0.473911	1132.462122
min	0.499900	1.000000	0.846154	0.333333	3.000000
25%	2.563400	18.000000	4.440716	1.006079	787.000000
50%	3.534800	29.000000	5.229129	1.048780	1166.000000
75%	4.743250	37.000000	6.052381	1.099526	1725.000000
max	15.000100	52.000000	141.909091	34.066667	35682.000000

	AveOccup	Latitude	Longitude	Target
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704	2.068558
std	10.386050	2.135952	2.003532	1.153956
min	0.692308	32.540000	-124.350000	0.149990
25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010

## Check for missing values

In [4]:

```
print("\nMissing values in dataset:")
print(df.isnull().sum())
```

Missing values in dataset:

MedInc	0
HouseAge	0
AveRooms	0
AveBedrms	0
Population	0
AveOccup	0
Latitude	0
Longitude	0
Target	0

dtype: int64

## Splitting Data

In [5]:

```
X = df.drop(columns=['Target'])
y = df['Target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

## Feature Scaling

In [6]:

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## Initialize models

In [7]:

```
models = {
    "Linear Regression": LinearRegression(),
    "Decision Tree": DecisionTreeRegressor(random_state=42),
    "Random Forest": RandomForestRegressor(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(n_estimators=100, random_state=42),
    "SVR": SVR()
}
```

## Model Training and Evaluation

```
In [8]: results = {}
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    results[name] = {"MSE": mse, "MAE": mae, "R2 Score": r2}

print(f"{name} Results:")
print(f"MSE: {mse:.4f}, MAE: {mae:.4f}, R2 Score: {r2:.4f}\n")
```

Linear Regression Results:

MSE: 0.5559, MAE: 0.5332, R2 Score: 0.5758

Decision Tree Results:

MSE: 0.4940, MAE: 0.4539, R2 Score: 0.6230

Random Forest Results:

MSE: 0.2552, MAE: 0.3274, R2 Score: 0.8053

Gradient Boosting Results:

MSE: 0.2940, MAE: 0.3717, R2 Score: 0.7756

SVR Results:

MSE: 0.3570, MAE: 0.3986, R2 Score: 0.7276

## Convert results to DataFrame and display

```
In [10]: results_df = pd.DataFrame(results).T
print("Model Performance Comparison:")
print(results_df)
```


Model Performance Comparison:

	MSE	MAE	R2 Score
Linear Regression	0.555892	0.533200	0.575788
Decision Tree	0.493969	0.453904	0.623042
Random Forest	0.255170	0.327425	0.805275
Gradient Boosting	0.293999	0.371650	0.775643
SVR	0.357004	0.398599	0.727563

## Visualization

```
plt.figure(figsize=(10, 5)) sns.barplot(data=results_df, palette='viridis') plt.title("Model Performance Comparison") plt.ylabel("Error / Score") plt.xticks(rotation=45) plt.show()
```

In [12]:



In [ ]: